

Improving PPI Prediction with ESM-derived Features

戴佳文 齐奕婷 王安瑞

1. Feature Extraction

- Utilized **ESMC** for protein sequence embedding
 - **ESM Cambrian** focuses on creating representations of the underlying biology of proteins.
 - Old methods use biological annotations or only protein sequences
- Technical Details:
 - Input: Protein sequences from the `b4ppi` dataset
 - Model: ESM-C pre-trained transformer
 - Output: Embeddings with shape $[L + 2, 960]$, where L represents sequence length and 960 is the embedding dimension.(other 2 are [CLS], [EOS] token)

2. Data Processing


To enable downstream classification, we explored multiple pooling strategies to transform variable-length embeddings $[L, 960]$ into fixed-size representations $[960]$.

Initial Approach: average pooling/max pooling

- Applied **average pooling** along the sequence dimension $L+2$ to:
 - Handle variable-length inputs uniformly
 - Reduce computational complexity (the embedding matrix is too big)
 - Preserve channel-wise information
- Resulting compressed embedding: $[1, 960]$ per protein
- Lose important positional and structural information

Improved Approach: Masked Autoencoder (MAE)

- Length Standardization = 1502
 - Pad shorter sequences with zeros ([PAD] = 0)
 - Truncate longer sequences
 - 75% mask
- Problems:
 - data with small variances 0.0002-0.001 and small loss ----> multiply a scale_factor
 - mask padding positions ----> calculate the loss only for non-padding positions




Key Improvement:

- Added **padding position markers** for each protein sequence

```
return {  
    "seq": seq.clone(),                # (max_len, 960)  
    "padding_start": seq_len           # int  
}
```

- Modified masking to only operate on **non-padding regions**.
 - 50% mask



other modifications:

- We first concatenated the two features and then applied compression; because the data volume was too large, we truncated the length of the concatenated protein embeddings to 2000.
- Use the CLS token added by ESMC as the compressed embedding.

Downstream Supervised Learning Models

We evaluated multiple classification algorithms to determine the optimal approach for PPI prediction.

- **Dataset Characteristics:**

- Test1 dataset is a balanced dataset with 50% positive and 50% negative samples.
- Test2 dataset presents a realistic challenge with 9.09% positive and 90.91% negative samples, closely mimicking real-world PPI prediction scenarios

L2-cos Method

Similarity	Dataset	Best AUROC
L2	Test1	0.5547
L2	Test2	0.5718
Cosine	Test1	0.5330
Cosine	Test2	0.5676

- High similarity between proteins does not necessarily imply that they have a PPI.

Logistic Regression

1. **As baseline:** Tests linear separability of features
2. **Computational efficiency:**
 - Fast training
 - Easy hyperparameter tuning (e.g., regularization strength C)
3. **Explainability:** Direct biological interpretation of feature contributions

Method	Dataset	AUROC
L2 Similarity	Test1	0.5547
L2 Similarity	Test2	0.5718
Logistic Regression	Test1	0.6182
Logistic Regression	Test2	0.6500

Support Vector Machine (SVM)

- Finds optimal decision boundary with regularization
- Use different kernels to capture non-linearities:

```
"svm__kernel" : ['linear', 'poly', 'rbf', 'sigmoid']  
# Handles linear/nonlinear cases```
```

- Effective in feature-rich, sample-limited scenarios

Method	Dataset	AUROC
Logistic Regression	Test1	0.6182
Logistic Regression	Test2	0.6500
SVM	Test1	0.6285
SVM	Test2	0.6201

XGBoost

Nonlinear feature interactions:

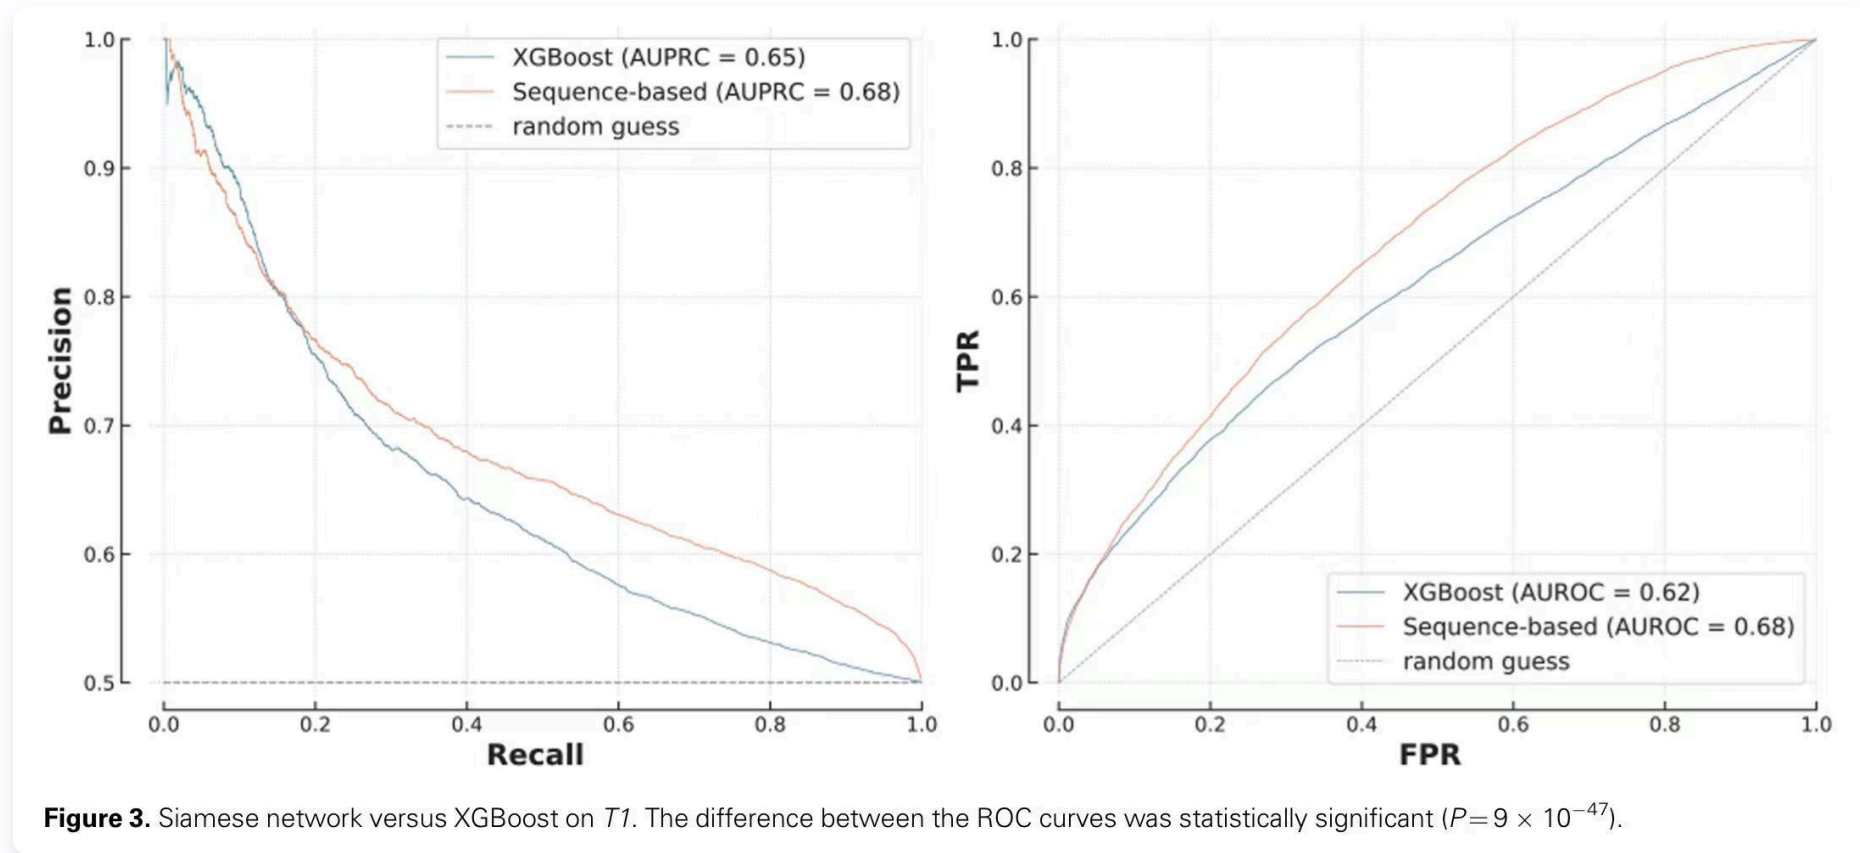
- Discovers complex patterns in 960-D space

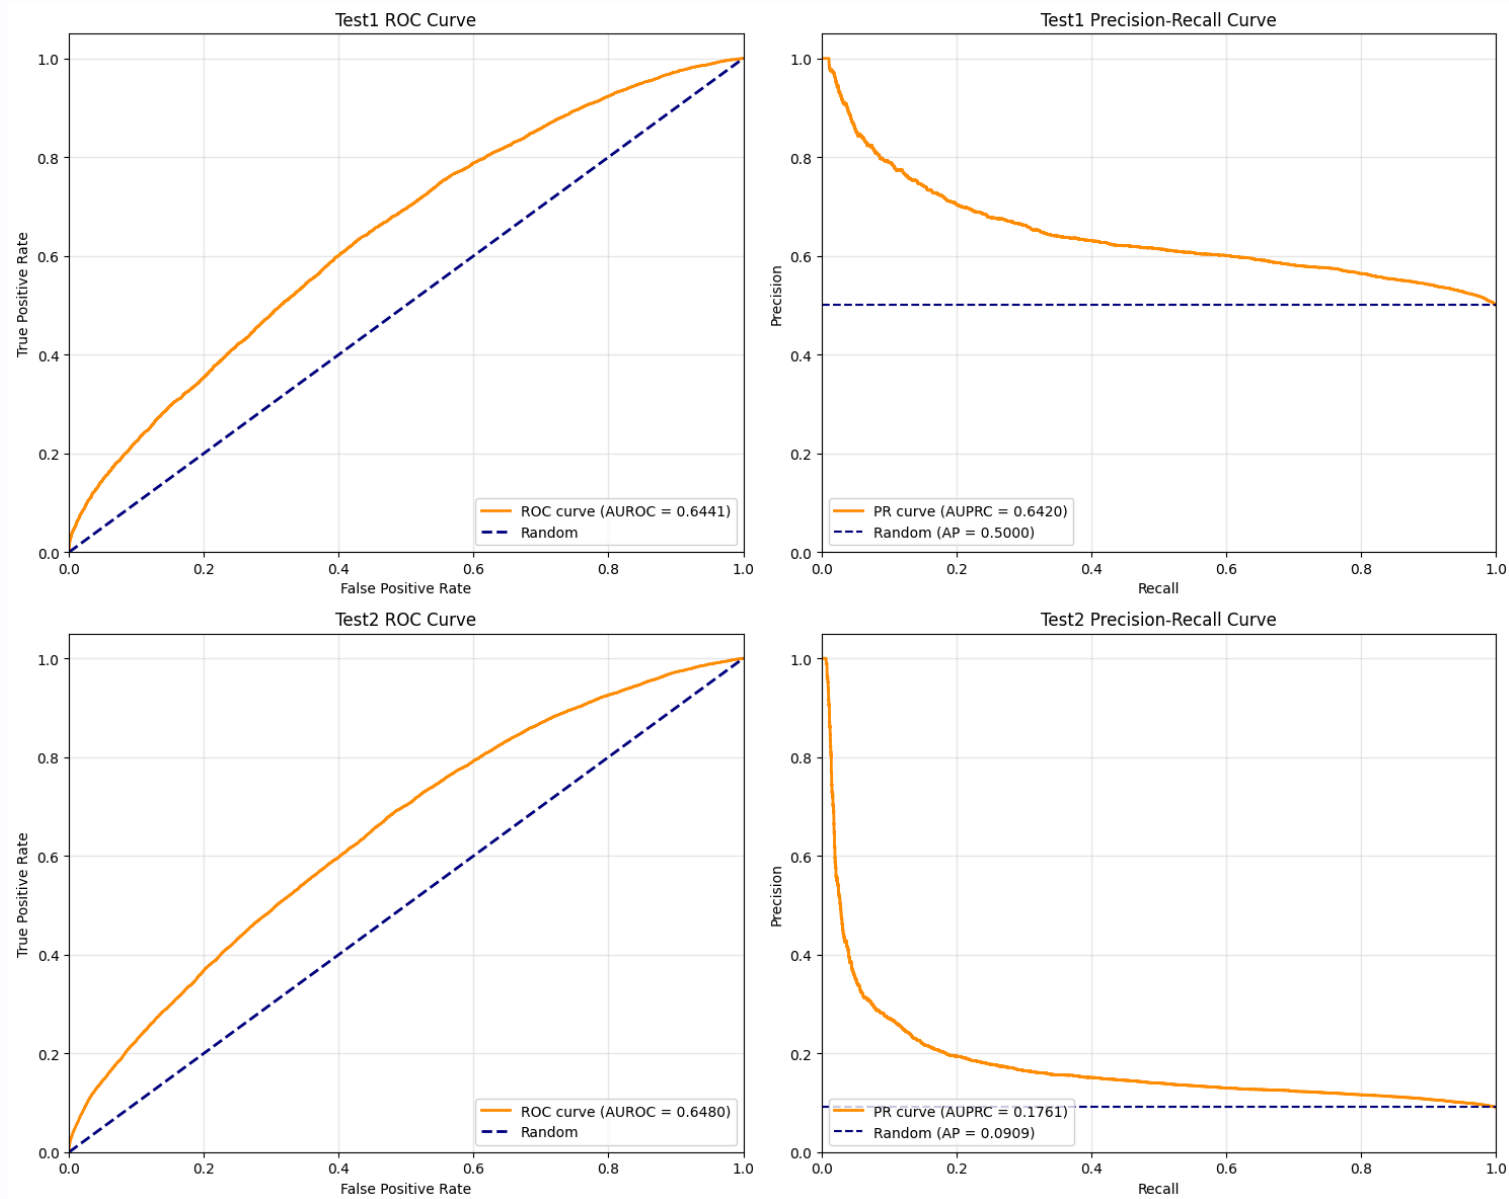
Method	Dataset	AUROC
Logistic Regression	Test1	0.6182
Logistic Regression	Test2	0.6500
XGBoost	Test1	0.6441
XGBoost	Test2	0.6480

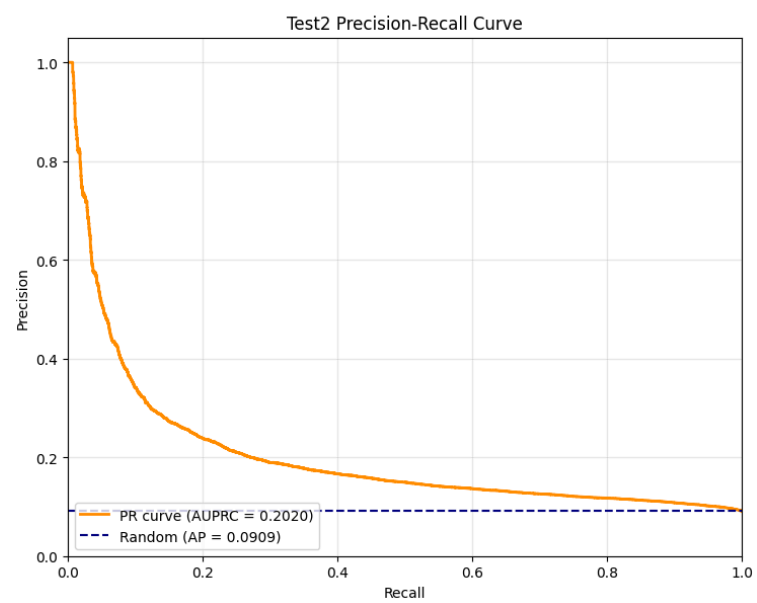
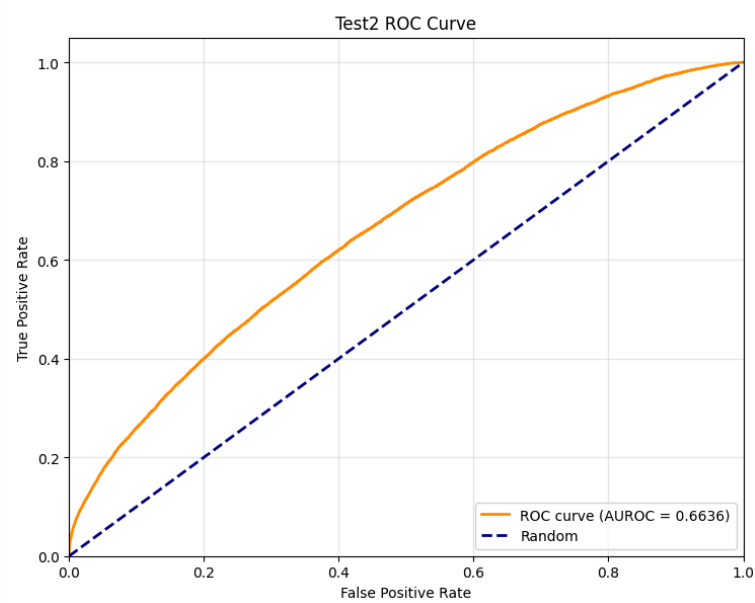

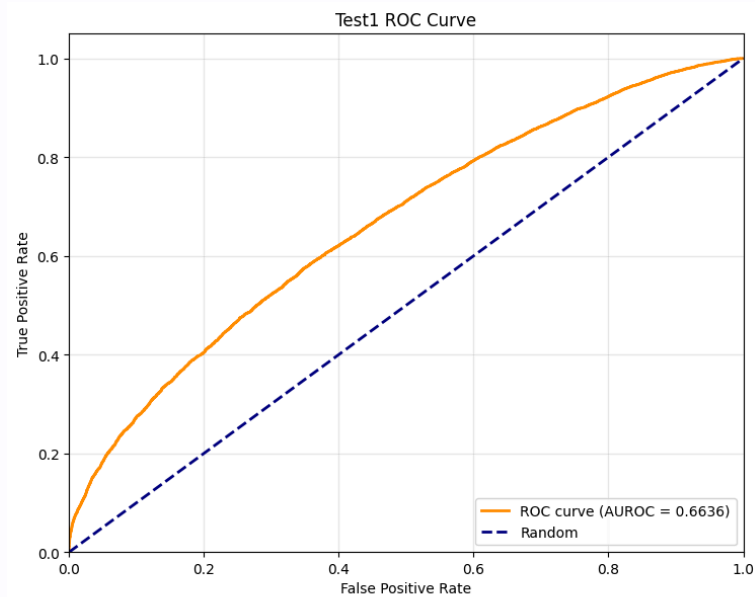
Multi-Layer Perceptron (MLP)

1. The encoder of MAE is essentially a multi-layer fully connected structure, and using MLP as the downstream classification layer can naturally connect. The encoder and MLP can be merged into a large end-to-end model.
2. Captures higher-order interactions than XGBoost
3. A unified backpropagation pipeline can be useful to fine-tune both the MAE encoder and MLP

Reference: "Pitfalls of machine learning models for protein–protein interaction networks"







Final Comparison Table (Test1 Only)

Method	Test Set	AUROC	AUPRC
Sequence-based (RNN, from literature)	Test1	0.68	0.68
XGBoost (from literature)	Test1	0.62	0.65
XGBoost + AvgPooling	Test1	0.6441	0.6447
	Test2	0.6480	0.1761
XGBoost + Masked Autoencoder (MAE)	Test1	0.6466	0.6771
	Test2	0.6636	0.2020

Conclusion

- **Competitive Results:** Our XGBoost + MAE approach outperforms literature XGBoost baselines, but achieves slightly lower AUROC than sequence-based RNN models from literature.
- **XGBoost Superiority:** XGBoost demonstrates the best performance among traditional ML classifiers, effectively capturing nonlinear feature interactions

Future Directions

- End-to-end fine-tuning of MAE encoder with downstream classifiers(MLP)
- Exploration of more sophisticated attention mechanisms